

# Getting Help

---

Main Instructor: **Azzam Mourad** <[azzam.mourad@nyu.edu](mailto:azzam.mourad@nyu.edu)>

- Office hours: online and in-person, MW: 11:15AM-12:30 pm or by appointment
- Zoom Link is provided in the syllabus

Teaching Assistant: **Khalid Mengal** <[kqm1@nyu.edu](mailto:kqm1@nyu.edu)>

- Office Hours: W: 2:00PM-3:00PM

# Course Overview

Computer Systems Organization

# Course Perspective

---

Most Systems Courses are Builder-Centric

- Computer Architecture
  - Design pipelined processor in Verilog
- Operating Systems
  - Implement large portions of operating system
- Compilers
  - Write compiler for simple language
- Networking
  - Implement and simulate network protocols

# Course Perspective (Cont.)

---

This course is **programmer-centric**

- Understanding of underlying system makes a more effective programmer
- Bring out the hidden hacker in everyone

# Textbooks

---

Randal E. Bryant and David R. O'Hallaron,

“**Computer Systems: A Programmer's Perspective**, 3rd Edition”,  
Prentice Hall

Brian Kernighan and Dennis Ritchie,

“**The C Programming Language**, 2nd Edition”, Prentice Hall

# Course Components

---

## Lectures

- Higher level concepts

## Programming Assignments/Labs (4)

- The heart of the course
- Provide in-depth understanding of some aspect of systems

## In-class Quizzes (2)

One midterm

One final exam

# Course Syllabus

---

- C Programming
- Data representation and manipulation (bit, int, float ...)
- Assembly and Program Representation
- Memory hierarchy
- Program optimization and parallelism
- Virtual memory
- Linking

# Abstraction Is Good But Don't Forget Reality

---

Most CS Courses emphasize abstraction

Goal of CSO:

- Help you understand how computers work and build bug-free/efficient software and system programs

Why is it important?

- Fundamental computer science
- Become better programmer
  - Avoid bugs
  - Write fast code
  - Write secure code
  - Write low-level/system code
- Prepare for later “systems” classes in CS
  - Operating Systems, Networking, Computer Architecture, Distributed Systems, Compilers



So, why CSO?

# Reason #1: Understanding Internal Representations is Important

---

$x^2 \geq 0$ ?

- $40000 * 40000 = ?$
- $50000 * 50000 = ?$


$(x + y) + z = x + (y + z)$ ?

- $(1e20 - 1e20) + 3.14 = 3.14$
- $1e20 - (1e20 + 3.14) = ?$

**Demo!**

# Ariane 5 Rocket

---



# Reason #2: Knowing Assembly is Important

---

No need to program in assembly

Knowledge of assembly helps one understand machine-level execution

- Creating/fighting malware
- Debugging
- Writing system software (e.g. compilers , OS)

# Reason #3: Speed is Very Important

---

Compute Intensive Applications are Everywhere



**Scalable Web Applications**




**Scalable Mobile Applications**




**Image Processing**



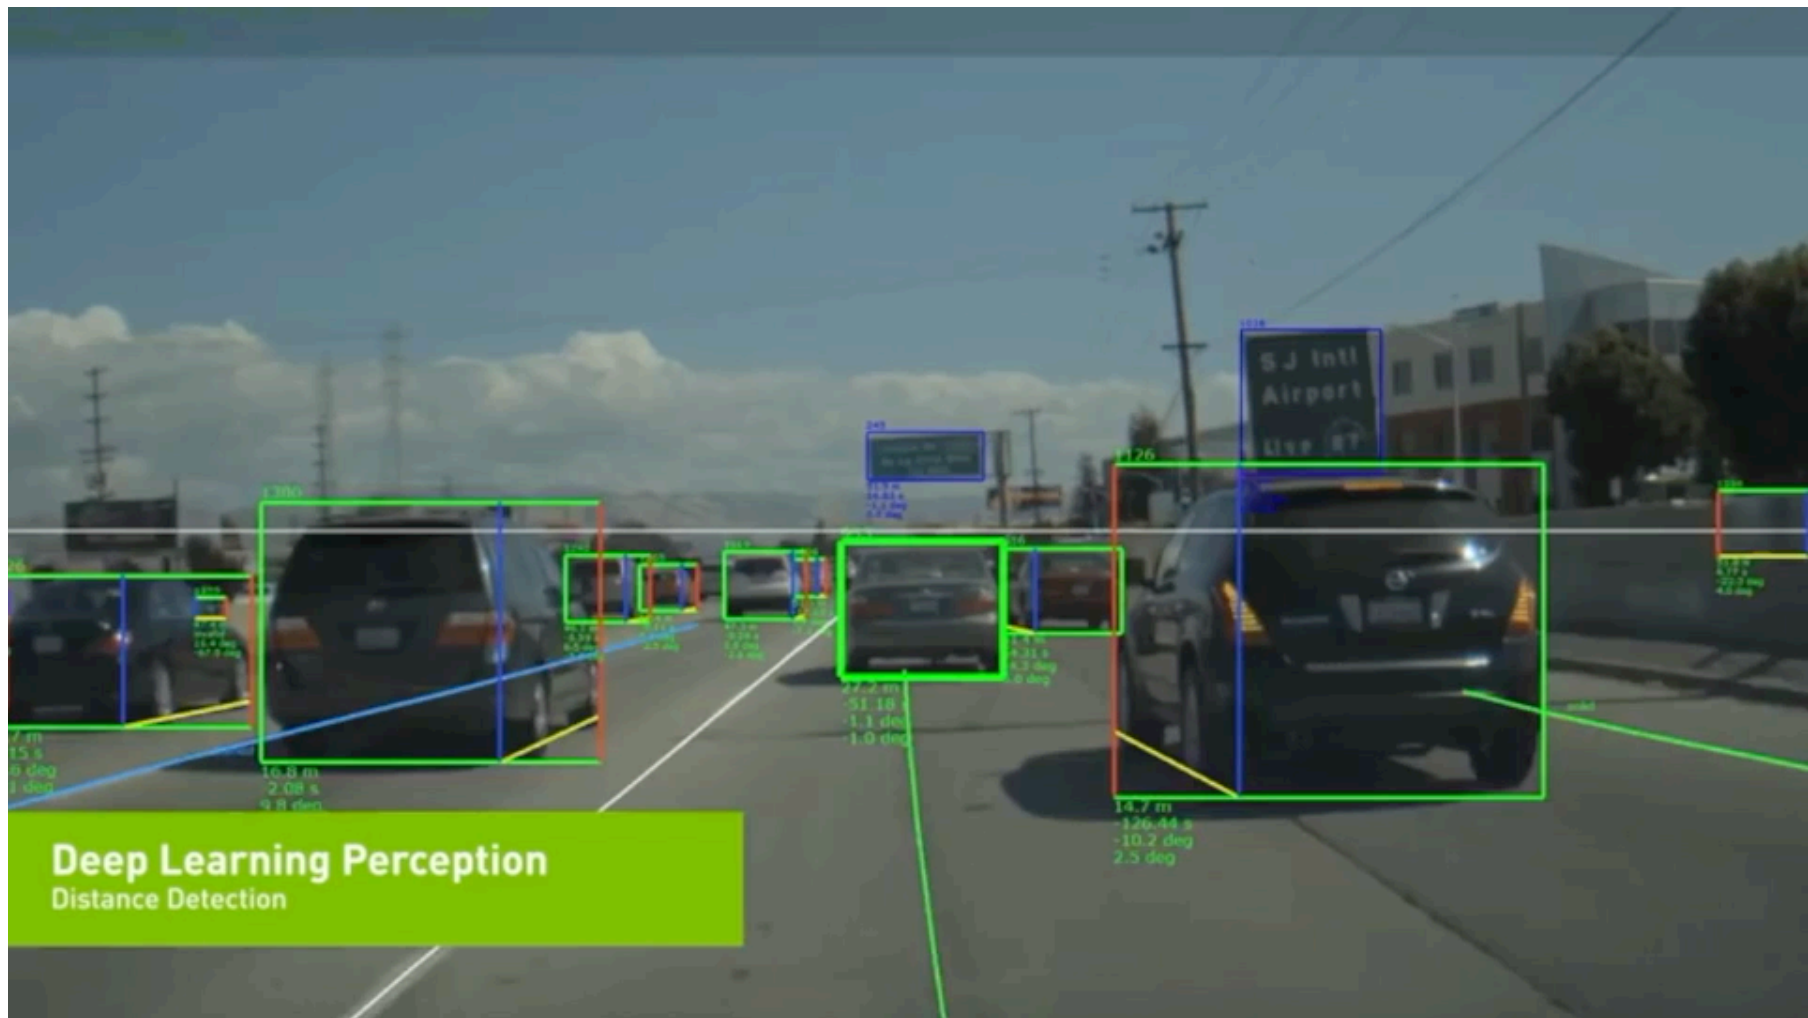
**Bioinformatics**



**Deep Learning**



**Scientific Computing**



**Driverless cars**

# Code Optimization in C

---



**Elon Musk**  @elonmusk · Feb 2

Our NN is initially in Python for rapid iteration, then converted to C++/C/raw metal driver code for speed (important!).



C is a good balance between high speed and productivity

# Memory System Performance Example

Is the right code faster?

```
for (i = 0; i < 2048; i++)  
  for (j = 0; j < 2048; j++)  
    dst[i][j] = src[i][j];
```

```
for (j = 0; j < 2048; j++)  
  for (i = 0; i < 2048; i++)  
    dst[i][j] = src[i][j];
```

21 times slower

Performance depends on access patterns



## Reason #4: Important for Security


---

Understanding how computers work is important for securing your code/hacking codes

# Access to Kernel Memory

---

Memory



# Memory Referencing Errors

---

C/C++ let programmers make memory errors

- Out of bounds array references
- Invalid pointer values
- Double free, use after free

Errors can lead to nasty bugs

- Corrupt program objects
- Effect of bug observed long after the corruption

# Memory Referencing Bug Example

---

```
double fun(int i)
{
    double d[1] = {3.14};
    int a[2];
    a[i] = 1073741824; /* Possibly out of bounds */
    return d[0];
}
```

fun(0)	→	3.14
fun(1)	→	3.14
fun(2)	→	?
fun(3)	→	?
fun(4)	→	?

# Code Security Example

---


```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy data_amount bytes from kernel region to user buffer */
int copy_from_kernel(void *user_buffer, int data_amount)
{
    /* Byte count len is minimum of buffer size and maxlen */
    int len = data_amount > KSIZE ? KSIZE : data_amount;
    memcpy(user_buffer, kbuf, len);
    ...
}
```

- Similar to code found in FreeBSD's implementation of getpeername (get name of connected peer socket)
- There are legions of smart people trying to find weaknesses in programs

# Code Security Example

```
/* Kernel memory region holding user-accessible data */  
#define KSIZE 1024  
char kbuf[KSIZE];  
  
/* Copy data_amount bytes from kernel region to user buffer */  
int copy_from_kernel(void *user_buffer, int data_amount)  
{  
    /* Byte count len is minimum of buffer size and maxlen */  
    int len = data_amount > KSIZE ? KSIZE : data_amount;  
    memcpy(user_buffer, kbuf, len);  
    ...  
}
```



# Typical Usage

---

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy data_amount bytes from kernel region to user buffer */
int copy_from_kernel(void *user_buffer, int data_amount)
{
    /* Byte count len is minimum of buffer size and maxlen */
    int len = data_amount > KSIZE ? KSIZE : data_amount;
    memcpy(user_buffer, kbuf, len);
    ...
}
```

```
#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, MSIZE);
    printf("%s\n", mybuf);
}
```

# Malicious Usage

---

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy data_amount bytes from kernel region to user buffer */
int copy_from_kernel(void *user_buffer, int data_amount)
{
    /* Byte count len is minimum of buffer size and maxlen */
    int len = data_amount > KSIZE ? KSIZE : data_amount;
    memcpy(user_buffer, kbuf, len);
    ...
}
```

```
#define MSIZE 528


void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, -MSIZE);
    . . .
}
```



# Malicious Usage

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy data_amount bytes from kernel region to user buffer */
int copy_from_kernel(void *user_buffer, int data_amount)
{
    /* Byte count len is minimum of buffer size and maxlen */
    int len = data_amount > KSIZE ? KSIZE : data_amount;
    memcpy(user_buffer, kbuf, len);
    ...
}
```



# Policies: Assignments (Labs)

---

You must work alone on all assignments

- Post all questions on the forum
- You are encouraged to answer others' questions, but refrain from explicitly giving away solutions

Hand-ins

- Assignments due at 11:55pm on the due date
- Late submissions: 10% deducted each late day (maximum 3 days)
- Two grace days
- Zero score if a lab is handed in > 3 days late

# UNIX Lab Environment

---

Use official class VM image

- Download (free) virtualbox for Windows/Linux
- Download VM appliance from course web page

**Your assignments must work on this environment!**

# Cheating

---

## What is cheating?

- Sharing code: by copying, looking at others' files
- Coaching: helping your friend write a lab step by step
- Copying code from a previous course or from elsewhere

## Penalty for cheating:

- Immediate removal from course with failing grade
- Permanent mark on your record