

Machine-Level Programming IV: Data

Slides adapted from the CMU version of the course (thanks to Randal E. Bryant and David R. O'Hallaron)

Today

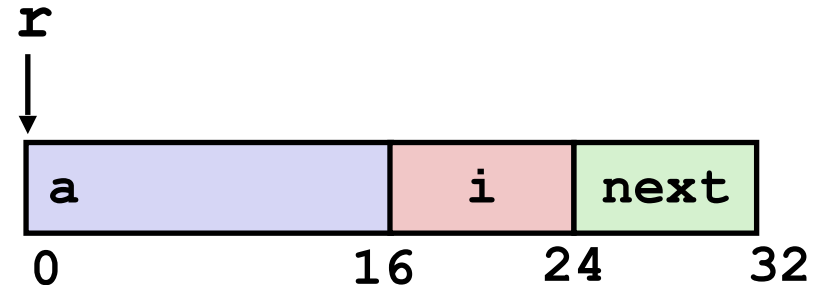
■ Structures

- Allocation
- Access
- Alignment

Structure Representation

```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```

Composite data type (multiple members)



- Structure represented as block of memory
 - Big enough to hold all of the fields (members)
- Fields ordered according to declaration
 - Even if another ordering could yield a more compact representation
- Compiler determines overall size + positions of fields
 - Machine-level program has no understanding of the structures in the source code

Side Note (&r->a[idx])

```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```

```
int *get_ap  
(struct rec *r, size_t idx)  
{  
    return &r->a[idx];  
}
```

■ Use of -> in structures

- Shortcut: Instead of writing (*r).a[idx] we can write r->a[idx]

■ What does &r->a[idx] mean?

- &(r->a[idx]) or (&r)->a[idx] ?

■ Answer

- It means &(r->a[idx]) because -> has a high precedence

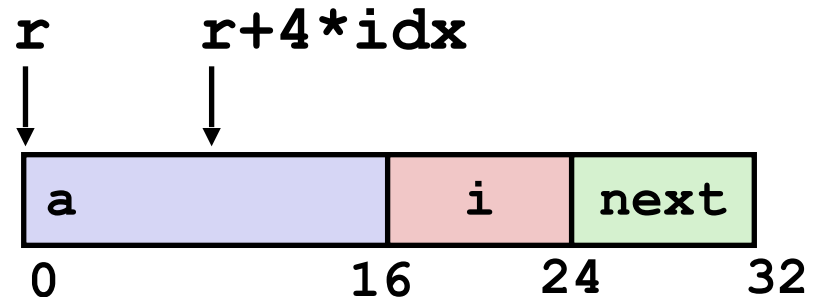
Generating Pointer to Structure Member

```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```

```
int *get_ap  
(struct rec *r, size_t idx)  
{  
    return &r->a[idx];  
}
```

■ Generating Pointer to Array Element

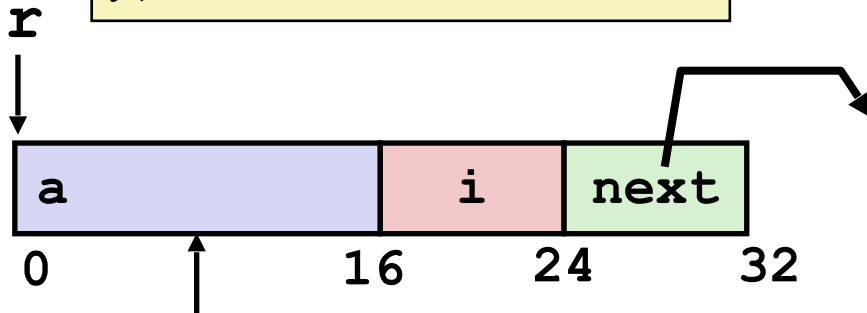
- Offset of each structure member determined at compile time
- Compute as $r + 4 * idx$



```
# r in %rdi, idx in %rsi  
leaq  (%rdi,%rsi,4), %rax  
ret
```

Following Linked List

```
struct rec {
    int a[4];
    size_t i;
    struct rec *next;
};
```



Element i

```
void set_val
(struct rec *r, int val)
{
    while (r) {
        size_t c = r->i;
        r->a[c] = val;
        r = r->next;
    }
}
```

Register	Value
%rdi	r
%rsi	val

```
.L11:                                # loop:
    movq    16(%rdi), %rax            # c = *(r+16)
    movl    %esi, (%rdi,%rax,4)      # *(r+4*c) = val
    movq    24(%rdi), %rdi           # r = *(r+24)
    testq   %rdi, %rdi               # Test r
    jne     .L11                     # if !=0 goto loop
```

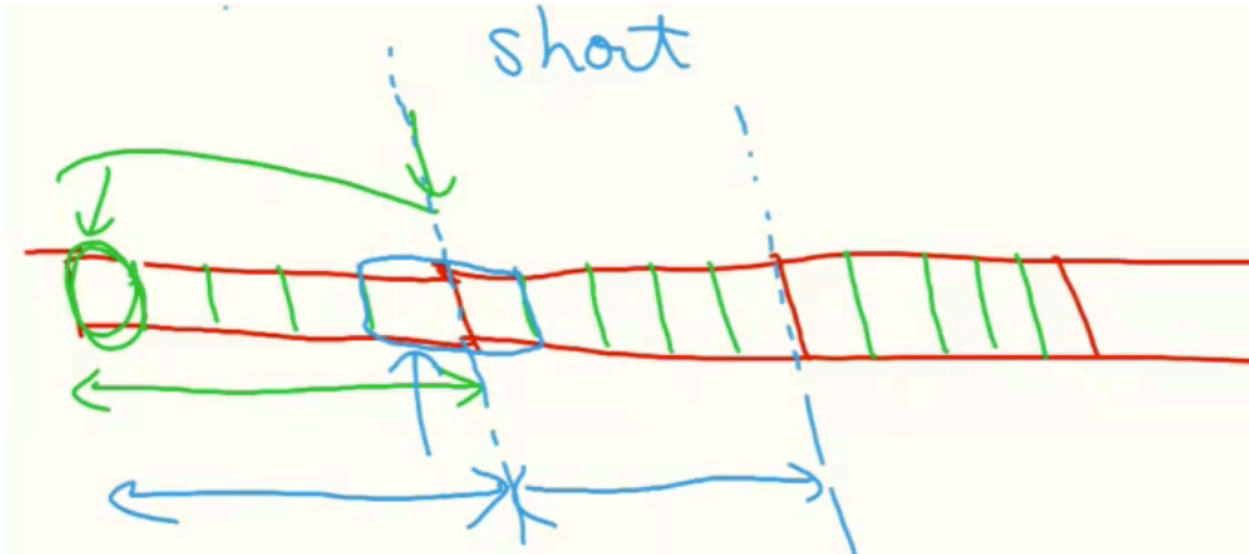
Today

■ Structures

- Allocation
- Access
- Alignment

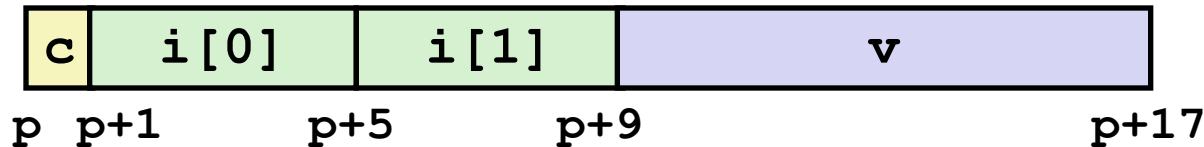
Structures & Alignment

■ Why do we Need Alignment?



Structures & Alignment

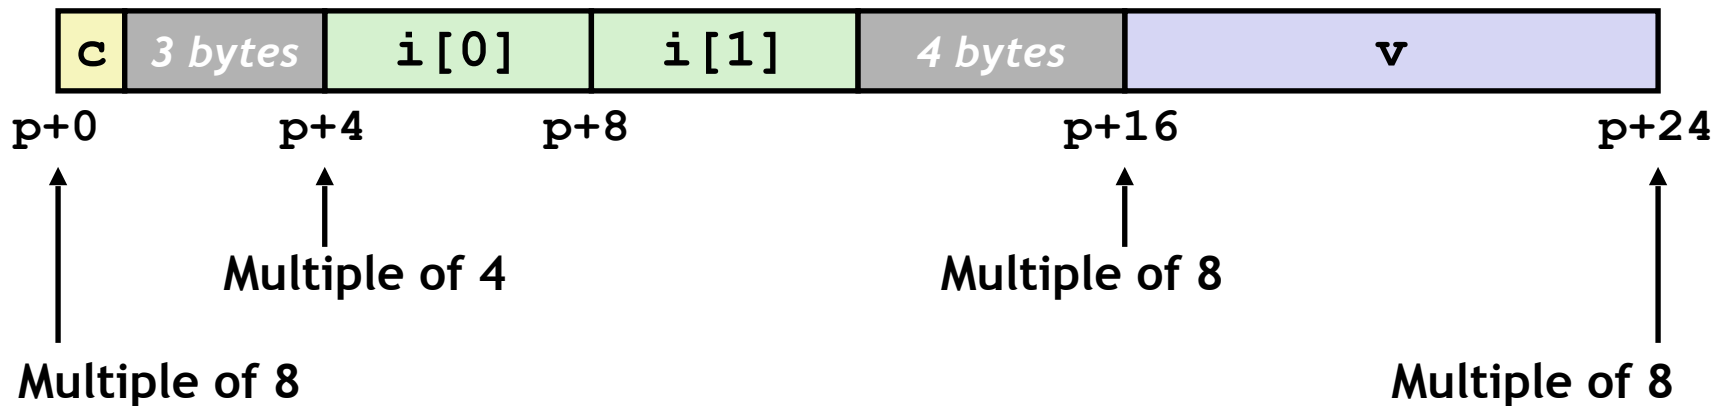
■ Unaligned Data



```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

■ Aligned Data

- If a primitive data type requires K bytes
- Then, the address must be multiple of K



Alignment Principles

■ Aligned Data

- Primitive data type requires K bytes
- Address must be multiple of K
- Required on some machines; advised on x86-64

■ Motivation for Aligning Data

- Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
 - Inefficient to load or store datum that spans quad word boundaries

■ Compiler

- Inserts gaps in structure to ensure correct alignment of fields

Satisfying Alignment with Structures

■ Within structure

- Must satisfy each element's alignment requirement

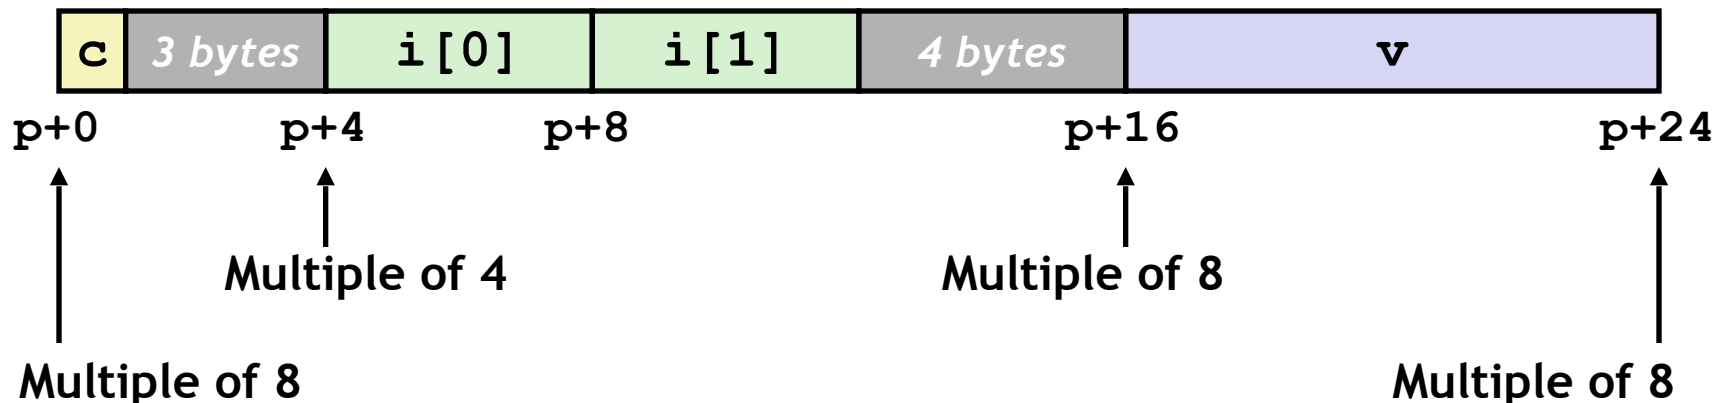
■ Overall structure placement

- Each structure has alignment requirement K
 - K = Largest alignment of any element
- Initial address & structure length must be multiples of K

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

■ Example:

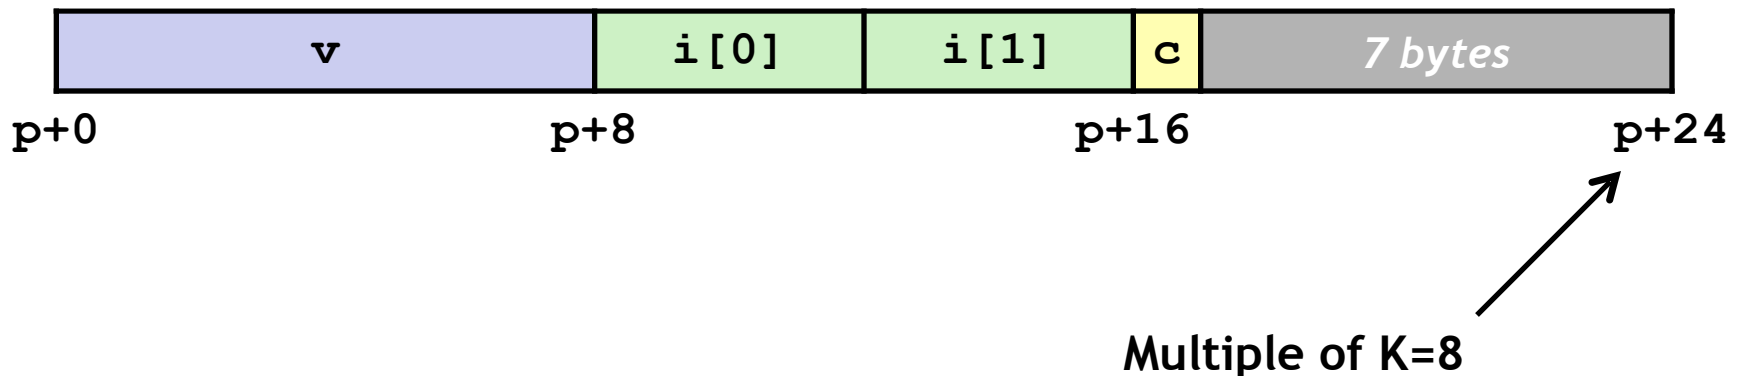
- $K = 8$, due to `double` element



Example 2 (Reorder Struct Fields)

- Assuming the order is different
- Now padding bytes are added to the end
- Why?
 - The largest alignment requirement is K
 - Overall structure size must be multiple of K

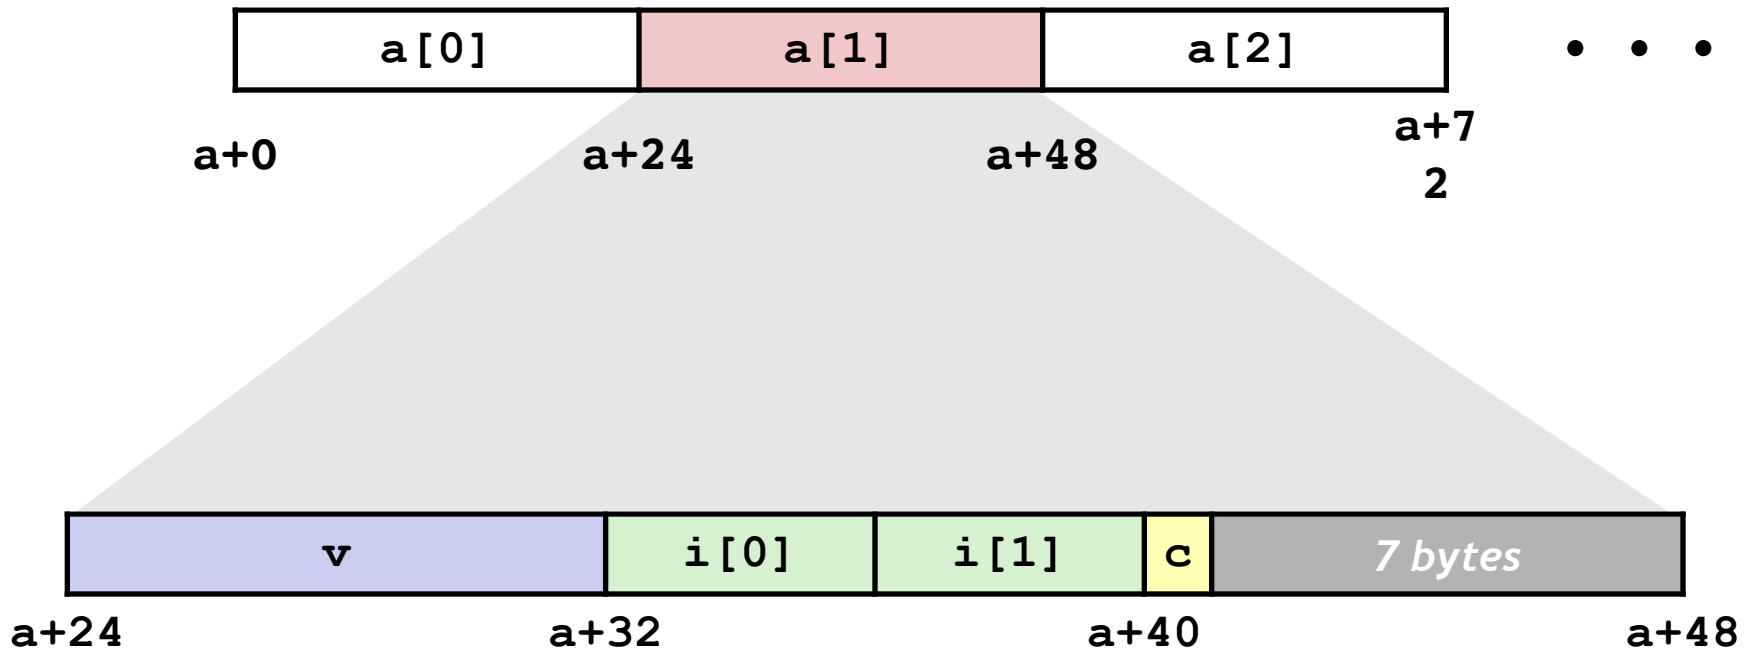
```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} *p;
```



Arrays of Structures

- Overall structure length multiple of K
- Satisfy alignment requirement for every element

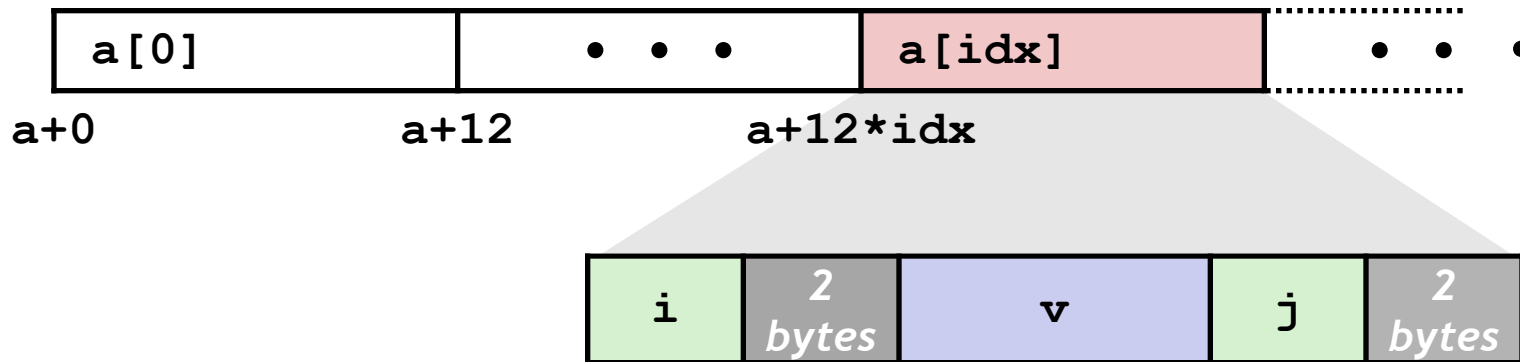
```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```



Accessing Array Elements

```
struct S3 {  
    short i;  
    float v;  
    short j;  
} a[10];
```

- Compute array offset $12 * \text{idx}$
- Element j is at offset 8 within structure
- Address: $a + 12 * \text{idx} + 8$



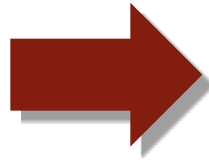
```
short get_j(int idx)  
{  
    return a[idx].j;  
}
```

```
# %rdi = idx  
leaq (%rdi,%rdi,2),%rax # 3*idx  
movzwl a+8(,%rax,4),%eax
```

Saving Space

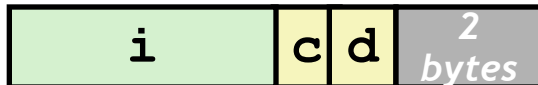
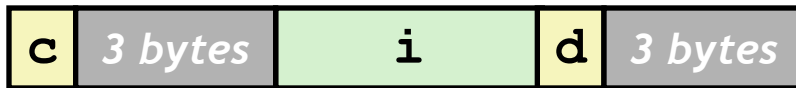
■ Put large data types first

```
struct S4 {  
    char c;  
    int i;  
    char d;  
} *p;
```



```
struct S5 {  
    int i;  
    char c;  
    char d;  
} *p;
```

■ Effect (K=4)



Practice

■ Reverse engineer the following code to C

```
struct S4 {  
    short v;  
    Struct S4 *p;  
};  
short f(struct S4 *ptr);
```

```
; short f(struct S4 *ptr)  
; ptr in %rdi  
  
    movl $1, %eax  
    jmp .L2  
.L3:  
    imulq (%rdi), %rax  
    movq 2(%rdi), %rdi  
.L2:  
    testq %rdi, %rdi  
    jne .L3  
    rep; ret
```


Practice

■ Reverse engineer the following code to C

```
struct S4 {  
    short v;  
    Struct S4 *p;  
};  
short f(struct S4 *ptr);
```

```
short f(struct S4 *ptr)  
{  
    short val = 1;  
    while (ptr)  
    {  
        val *= ptr->v;  
        ptr = ptr->p;  
    }  
    return val;  
}
```

```
; short f(struct S4 *ptr)  
; ptr in %rdi  
  
    movl $1, %eax  
    jmp .L2  
.L3:  
    imulq (%rdi), %rax  
    movq 2(%rdi), %rdi  
.L2:  
    testq %rdi, %rdi  
    jne .L3  
    rep; ret
```